

Screened at the door, ruled by a human, receipted on the way out

Opt-in per deployment (UNDERCROFT_ADMISSION=quarantine) — off, the write contract is byte-identical. On, this is the path every save takes.

A save arrives

MCP tool · REST /v1 · CLI · bulk ingest — handler stamps added_by: no writer picks its surface. Updates are re-stamped with the surface writing NOW — never the original writer's standing. Provenance claims (agent/channel/session) are HMAC-covered — **evidence, never a boundary**.

Trusted surface? Straight in.

UNDERCROFT_ADMIT_TRUSTED_SOURCES matches the stamped added_by only — a claim cannot admit itself. The bypass is the deployment's own declaration.

Tier 1 — deterministic. No model, no network; signal codes and byte offsets, never content.

A closed vocabulary, so the review queue reads as data. One signal per class, offsets into the candidate; classes record independent evidence.

imperative-instruction

tool-call-syntax

exfil-marker

encoded-blob

destination-anomaly

fixture-similarity

Windowed hash-embedder cosine against 18 committed attack fixtures — catches the variant that dodges every marker substring. Threshold pinned from both sides; measured 0/5,882 clean-corpus false positives, 18/18 fixtures trip their own screen.

rate-anomaly

The signal candidate bytes cannot carry: a declared per-writer budget (UNDERCROFT_ADMISSION_RATE=count/seconds), read from the vault's own history. Identity = agent claim, else the stamped surface among claim-less rows. Declared, never defaulted.

no signal

Tier 2 — a local model, advisory-only (UNDERCROFT_ADMISSION_LLM=advisory)

Consulted only for tier-1-clean candidates, only toward quarantine — flagged content is never shown to it: talking the classifier into "clean" bypasses nothing. A failed answer is a non-event, never a blocked write. The model is itself an injection target, steerable only in the safe direction.

Clean — lands where it aimed

Sealed, HMAC-tagged, chain-appended in the same transaction — the ordinary write path, untouched.

any signal → diverted, never rejected, never silently dropped

The quarantine wing — reserved, sealed, and outside every answer

The drawer keeps its verbatim content, sealed like any other; signal codes, offsets and the intended destination ride in metadata. Hard-excluded from retrieval except a reviewer's explicitly named scope — it can neither answer nor crowd. The wing refuses forged residents (a caller cannot aim a save here), quarantined drawers are not editable (the reviewer rules on exactly what the screen saw), and a flagged UPDATE diverts while the original drawer keeps its previous content.

allow — re-filed where it was headed

The human ruling IS the override, never re-screened — but refused when the destination moved since the text was queued (O224).

deny — destroyed, with a receipt

Rides the attested-forgetting path (C3.2): heads before/after, the tombstone interval, content fingerprints — verifiable by third parties.

Every ruling is a chain record with the verdict inside its tag's canonical — the review trail is as tamper-evident as the data. Detection is heuristic and says so: markerless poison passes, prose about injection can trip — which is exactly why a signal quarantines for a human and never rejects.