

Reading a date the writer already wrote

Four read-time declarations, and one signal that outranks all of them. Dates are read live, so a fix reaches existing vaults with no migration.

The marker beside the year outranks the calendar the caller declared

The declaration is a statement about a corpus. The marker is the writer's statement about THIS date. Reading it is evidence, so it wins.

Recognised, before the year, after it, or glued to it:

พ.ศ. · ค.ศ. · هـ · 民國 · 令和 · 平成 · 昭和

A bare year is a mention only where a marker names it.

so, with no declaration at all:

2568 is a quantity — พ.ศ. 2568 is 2025

Disagreeing markers settle nothing: the declaration stands.

Field order — four signals, strongest first

1 · declared on Locale

2 · demonstrated by the text itself
13/05 can only be day-first — the convention, stated by example

3 · implied by the language, per CLDR
Arabic is d/M/y in every territory. English splits: no order implied.

4 · else day-first — a documented default, not a confident reading

Ten calendars, two kinds of conversion

a renumbered year — arithmetic

Buddhist −543 · Minguo +1911 · 令和 +2018 · 平成 +1988
昭和 +1925 · 大正 +1911 · 明治 +1867

a different calendar — the whole date converts

Hijri, Umm al-Qura — Saudi CIVIL calendar printed on documents, deliberately not the tabular one, which is wrong by a day or two. Jalali turns at the vernal equinox, so no offset reaches it either.

A Japanese era is bounded, so its first and last years are PARTIAL

令和1年 = 1 May – 31 Dec 2019 — reading it as the whole of 2019 would claim four months that were 平成31年.

That is wrong rather than rounded, so the era carries a first and last day and the mention resolves to the period it actually names.

What is deliberately NOT read as evidence

Script is not a calendar. Thai script writes Gregorian dates constantly, and a corpus is not Buddhist because of the glyphs it uses.

A numeral system is not a calendar either: ๒๐๒๖ is an ordinary Gregorian 2026 typed in Thai digits. Reading the glyphs as an era claim resolved it to 1483 — a real bug, and the reason the rule is stated this bluntly.

A number's SIZE is not evidence of its kind either — so a range check is the wrong fix for the case below. The noun governing it IS evidence.

A wrong reading, and the signal that would fix it

م١٩٩٥ is how Arabic writes a year; م ١٥٠٠ how it writes a quantity. The difference is one space.

So م٢٥٠٠ على ارتفاع is an altitude in METRES, and the engine reads it as the year 2500.

A bare م counts when a YEAR NOUN governs the number OR the marker is glued. Only the second fires — and ارتفاع goes unread.

Two open gaps, recorded as gaps

The altitude above is not an impossibility — the governing noun is in the token stream. A year noun is already read as confirming evidence; a measurement noun pointing the other way is the same class of signal, and is simply not consulted yet.

Separately, in-text scanners exist for English and Arabic only. Greek and Russian have full paradigm coverage and still cannot read a date in their own text — that machinery is language-agnostic.