

Layers and trust boundaries

One palace per process. Tenancy is vaults, not palaces. Nothing leaves the machine unless you configure it to.

callers

agent over MCP studio
34 tools · local process

app over HTTP /v1
bearer token · REST

CLI
remember · mine · search

admin console /ui
served on every build

↓ bearer + optional per-vault assertion (HMAC, time-boxed)

engine — one Rust process

/v1 surface + MCP handler
drawers · search · refine · kg · verify
rotate · export · import · taxonomy

store
rank → gate → rerank → select
verifies every row's HMAC before use

knowledge graph
receipted facts, validity windows
written only by refine or explicit kg calls

optional local LLM client
Ollama or OpenAI-compatible, for refine
unset UNDERCROFT_LLM_URL ⇒ never called

optional, separate binary

orchestrator
tenant → vault routing, /t/*
own SQLite, sealed engine creds

remote vector backends
untrusted accelerators — sealed
payloads only, re-verified on read

↓ everything below this line is ciphertext

at rest — per-vault key, HKDF-derived from the master key

sealed artifacts
content · embeddings · ColBERT token
matrices · PQ codes, codebook, pages
each under a distinct AAD domain

integrity
per-row HMAC tag, verified on every
read; tamper-evident audit chain;
manifest anchor as rollback detector

never present for sealed vaults
no full-text index, no plaintext
derived data in clear on disk.
Search uses decrypt-once RAM caches.

Default posture: local-first and offline. No telemetry deps compiled in, no external API, deterministic embedder.

Each outward path is a deliberate opt-in: UNDERCROFT_LLM_URL, UNDERCROFT_OTLP_ENDPOINT, a remote index URL, UNDERCROFT_LLM_KEY.