

Write path — one drawer, and why it is safe

Content is stored verbatim. Only derived structure is added; the text itself is never rewritten.

1 • normalize

safety floor always: strip NUL and control chars, CRLF → LF

Code mode keeps every space

2 • derive, never alter

time_mentions · entities
content_date anchors relative dates
runs in Drawer::new — no path forgets

3 • embed

deterministic offline hash embedder
by default — no model, no network
computed before the lock is taken

4 • one SQLite transaction — all of it, or none of it

BEGIN IMMEDIATE

seal

XChaCha20-Poly1305 over content and embedding.

AAD = the drawer id, plus distinct domains /emb · /tok · /pq

tag

HMAC-SHA256 over
id || meta_json || sealed content
metadata is covered for free — a new meta field needs no schema change

chain

chain_append runs in the SAME transaction as the row.

so no committed write can exist without its audit entry

Bulk ingest: one transaction and one anchor per batch, not per drawer — measured ≈55× fewer fsyncs at 200 drawers.

5 • then, outside the transaction

fsync the manifest anchor through an atomic rename + directory sync. WAL, synchronous=FULL.

why the anchor lags on purpose

anchor behind the db at open → power loss, fast-forward
anchor ahead of the db at open → rollback, tamper alarm

The asymmetry is the detector

An anchor that never runs ahead of the database is what makes "the machine lost power" and "someone restored an older copy of this file" distinguishable at all. That is why the invariant is directional, not just "keep them in sync".